

機械学習勉強会2.0

GAN



GANの概要

GAN (Generative Adversarial Networks)

画像生成手法の一つ

progressive GAN



<https://arxiv.org/abs/1710.10196>


cycle GAN




<https://arxiv.org/abs/1703.10593>

GANの概要

生成器: 本物のデータと見分けが
付かないような偽のデータを作り出す

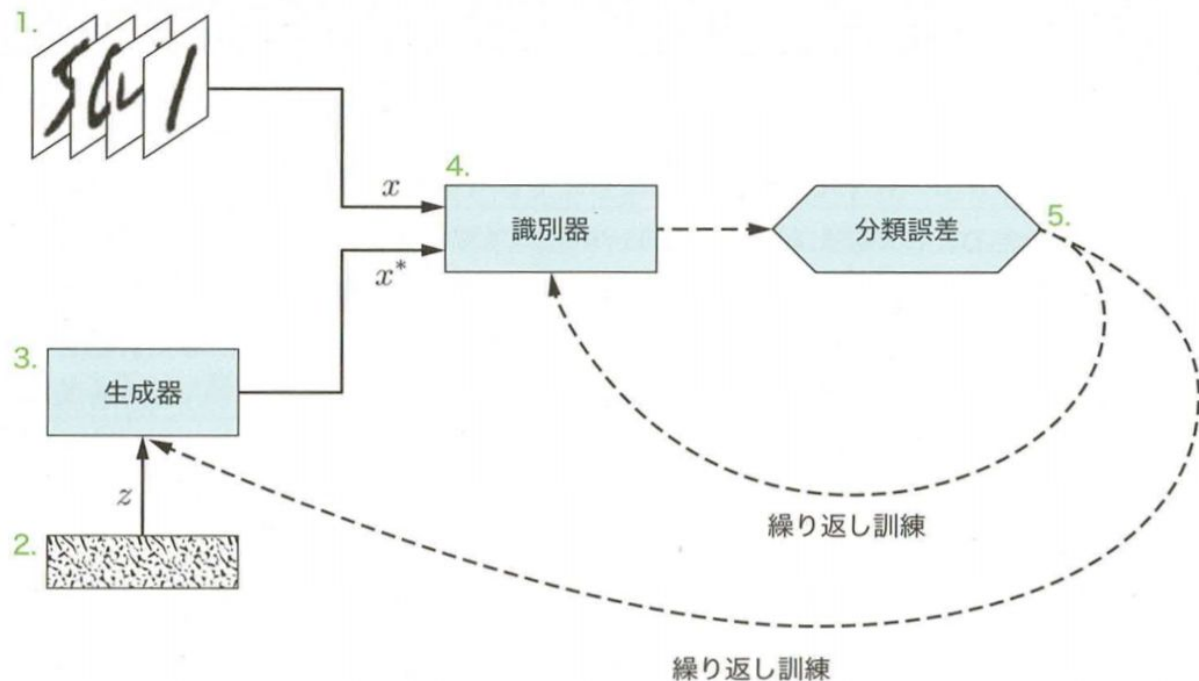


識別器: 本物のデータと、生成器が作り出した
偽のデータを区別する



敵対的: 生成器と識別器がゲームのような競争的な動作をする

GAN全体図




ナッシュ均衡に
落ち着くまで繰り返す。

DCGAN (Deep Convolutional GAN)

ConvNetとGANを組み合わせたもの

2016年にAlec Radford, Luke Metz, Soumith Chintalaらによって提案¹⁾



“A Guide to Convolution Arithmetic for Deep Learning”, <https://arxiv.org/abs/1603.07285>

1) <https://arxiv.org/abs/1511.06434>

GAN演習

演習

google colaboratoryでの実習を行う際には、下記リンクよりアクセス下さい。

https://colab.research.google.com/github/academeia/machine-learning-seminar_2020/blob/master/GAN/GAN.ipynb

ランタイムのタイプを変更することで、GPUによる計算が可能です。



初期設定、各種インポート

```
import tensorflow as tf
```

```
tf.__version__
```

```
'2.2.0'
```

```
# To generate GIFs  
!pip install imageio
```

```
import glob  
import imageio  
import matplotlib.pyplot as plt  
import numpy as np  
import os  
import PIL  
from tensorflow.keras import layers  
import time  
  
from IPython import display
```

データの取得、バッチ化


```
(train_images, train_labels), (_, _) = tf.keras.datasets.mnist.load_data()
```

```
train_images = train_images.reshape(train_images.shape[0],  
                                     28, 28, 1).astype('float32')  
train_images = (train_images - 127.5) / 127.5  
               # Normalize the images to [-1, 1]
```

```
BUFFER_SIZE = 60000  
BATCH_SIZE = 256
```

```
# Batch and shuffle the data  
train_dataset = \  
tf.data.Dataset.from_tensor_slices(train_images).shuffle(BUFFER_SIZE).batch(BATCH_SIZE)
```

生成器のモデル概略



生成器のモデル関数 (1)

```
def make_generator_model():
    model = tf.keras.Sequential()
    model.add(layers.Dense(7*7*256, use_bias=False, input_shape=(100,)))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())

    model.add(layers.Reshape((7, 7, 256)))
    assert model.output_shape == (None, 7, 7, 256)
                                   # Note: None is the batch size

    model.add(layers.Conv2DTranspose(128, (5, 5), strides=(1, 1),
                                     padding='same', use_bias=False))
    assert model.output_shape == (None, 7, 7, 128)
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())
```

生成器のモデル関数 (2)

```
model.add(layers.Conv2DTranspose(64, (5, 5), strides=(2, 2),
                                padding='same', use_bias=False))
assert model.output_shape == (None, 14, 14, 64)
model.add(layers.BatchNormalization())
model.add(layers.LeakyReLU())

model.add(layers.Conv2DTranspose(1, (5, 5), strides=(2, 2),
                                padding='same', use_bias=False,
                                activation='tanh'))
assert model.output_shape == (None, 28, 28, 1)

return model
```


生成器のモデル、初期画像

```
generator = make_generator_model()


noise = tf.random.normal([1, 100])
generated_image = generator(noise, training=False)

plt.imshow(generated_image[0, :, :, 0], cmap='gray')
```

<matplotlib.image.AxesImage at



識別器のモデル概略



識別器のモデル関数

```
def make_discriminator_model():  
    model = tf.keras.Sequential()  
    model.add(layers.Conv2D(64, (5, 5), strides=(2, 2), padding='same',  
                             input_shape=[28, 28, 1]))  
  
    model.add(layers.LeakyReLU())  
    model.add(layers.Dropout(0.3))  
  
    model.add(layers.Conv2D(128, (5, 5), strides=(2, 2), padding='same'))  
    model.add(layers.LeakyReLU())  
    model.add(layers.Dropout(0.3))  
  
    model.add(layers.Flatten())  
    model.add(layers.Dense(1))  
  
    return model
```


識別器、交差エントロピー

```
discriminator = make_discriminator_model()  
decision = discriminator(generated_image)  
print (decision)
```

```
tf.Tensor([[ -0.00156255]], shape=(1, 1), dtype=float32)
```

```
# This method returns a helper function to compute cross entropy loss  
cross_entropy = tf.keras.losses.BinaryCrossentropy(from_logits=True)
```

交差エントロピー

$$H(p, q) = - \sum_x p(x) \log q(x)$$

p : 正解の確率分布、 q : 予測の確率分布

損失関数、最適化アルゴリズム

```
def discriminator_loss(real_output, fake_output):  
    real_loss = cross_entropy(tf.ones_like(real_output), real_output)  
    fake_loss = cross_entropy(tf.zeros_like(fake_output), fake_output)  
    total_loss = real_loss + fake_loss  
    return total_loss
```

```
def generator_loss(fake_output):  
    return cross_entropy(tf.ones_like(fake_output), fake_output)
```

```
generator_optimizer = tf.keras.optimizers.Adam(1e-4)  
discriminator_optimizer = tf.keras.optimizers.Adam(1e-4)
```

$$H(p, q) = - \sum_x p(x) \log q(x)$$

チェックポイント、各種パラメーター

```
checkpoint_dir = './training_checkpoints'  
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt")  
checkpoint = tf.train.Checkpoint(generator_optimizer=generator_optimizer,  
  
discriminator_optimizer=discriminator_optimizer,  
                                generator=generator,  
                                discriminator=discriminator)
```

```
EPOCHS = 50  
noise_dim = 100  
num_examples_to_generate = 16  
  
# We will reuse this seed overtime (so it's easier)  
# to visualize progress in the animated GIF)  
seed = tf.random.normal([num_examples_to_generate, noise_dim])
```

訓練の各ステップ (1)

```
# Notice the use of `tf.function`  
# This annotation causes the function to be "compiled".  
@tf.function  
def train_step(images):  
    noise = tf.random.normal([BATCH_SIZE, noise_dim])  
  
    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:  
        generated_images = generator(noise, training=True)  
  
        real_output = discriminator(images, training=True)  
        fake_output = discriminator(generated_images, training=True)  
  
        gen_loss = generator_loss(fake_output)  
        disc_loss = discriminator_loss(real_output, fake_output)
```

訓練の各ステップ (2)

```
gradients_of_generator = \  
    gen_tape.gradient(gen_loss, generator.trainable_variables)  
gradients_of_discriminator = \  
    disc_tape.gradient(disc_loss, discriminator.trainable_variables)  
  
generator_optimizer.apply_gradients(zip(gradients_of_generator,  
                                         generator.trainable_variables))  
discriminator_optimizer.apply_gradients(zip(gradients_of_discriminator,  
discriminator.trainable_variables))
```

訓練用関数 (1)

```
def train(dataset, epochs):  
    for epoch in range(epochs):  
        start = time.time()  
  
        for image_batch in dataset:  
            train_step(image_batch)  
  
        # Produce images for the GIF as we go  
        display.clear_output(wait=True)  
        generate_and_save_images(generator, epoch + 1, seed)
```

画像の生成と保存

```
def generate_and_save_images(model, epoch, test_input):  
    # Notice `training` is set to False.  
    # This is so all layers run in inference mode (batchnorm).  
    predictions = model(test_input, training=False)  
  
    fig = plt.figure(figsize=(4,4))  
  
    for i in range(predictions.shape[0]):  
        plt.subplot(4, 4, i+1)  
        plt.imshow(predictions[i, :, :, 0] * 127.5 + 127.5, cmap='gray')  
        plt.axis('off')  
  
    plt.savefig('image_at_epoch_{:04d}.png'.format(epoch))  
    plt.show()
```

訓練

```
train(train_dataset,  
      EPOCHS)
```



Time for epoch 10 is 711.0001904964447 sec

```
checkpoint.restore(tf.train.latest_checkpoint(checkpoint_dir)  
)
```